

Markdown Example

Your Name

January ??, 2024

Contents

Instructions	1
R Markdown and Reproducible Research	1
Advantages of R scripts: Reproducible Analyses	2
Advantages of R Markdown: Simple markup and LaTeX formulas	2
Installing R and R Studio	3
Installing and loading packages	3
Reading data in a text file	4
Visualizing frequencies	8
Barchart of frequencies	8
Get rid of ‘Total’	9
Questions	16

Last updated on: January 09, 2024 at 19:33

Instructions

After installing R and RStudio using these instructions start RStudio in the project you created for this course.

Use the menus to create a new R script with *File > New File > R Script*.

Copy this file into the R Script and save it with *Control+S* (*Command+S* in macOS) with the name ‘markdown_sample.R’

Execute this file manually, line by line, using *Control+Enter* in Windows or *Command+Enter* on macOS. As you go through the file, make appropriate changes: e.g. fill in your name on line 3 above.

This will install a number of packages and download a data file.

After executing the file line by line you should be able to ‘render’ it to create a stand-alone HTML file by pressing *Control+Shift+K* (*Command+Shift+K* on macOS).

R Markdown and Reproducible Research

This is an example of output from a “.R script with R Markdown”. If you run this script in R, all lines that begin with ‘#’ are treated as comments. If you run it with ‘knitr’ – which just involves typing *Control-Shift-K* simultaneously when the file is the active file in R Studio – all the text in lines that begin with “# ” (hashtag apostrophe space) are processed as R Markdown code that can be used to produce a polished document in a number of standard formats: pdf, HTML, RTF, beamer presentations, etc.

Note that online help for R Markdown describes syntax for “.Rmd” files. They are the same as “.R scripts with R Markdown” files minus the “#’ ” and plus *code chunk delimiters*. In .R scripts, the code chunk delimiters are not necessary, which makes .R scripts easier to use. There are legitimate reasons to prefer “.Rmd” files but for uniformity when collaborating it is suggested that you use “.R scripts with R Markdown” in this course.

Almost all the work in this course will be done and submitted this way.

Advantages of R scripts: Reproducible Analyses

One of the great strengths of Markdown is that it allows you to do **reproducible** analyses and reports. Another researcher can use your R script on the same data and get the same results. They can verify the exact steps you took and can easily test how results would be affected by modifying the analysis.

You can collaborate with others much more easily knowing that if you send a script and data files to a collaborator they will get the same output (provided they have installed the same packages).

An important practice for reproducible research is that the analyst should never modify raw data. For example, if you find some errors in a spreadsheet sent by a client or downloaded from the internet, it is very tempting to manually correct the errors in the spreadsheet.

However, when you receive or download a new version of the spreadsheet you would need to apply the corrections the same way. If you did them manually you will not be able to guarantee consistency in your corrections.

To ensure reproducibility, the corrections should be done with code in the R script. This will usually involve the use of a very powerful tool: regular expressions, the subject of this xkcd cartoon.

With regular expressions, data corrections can usually be made in a way that does not introduce errors when applied to a future corrected version of the spreadsheet.

Advantages of R Markdown: Simple markup and LaTeX formulas

As illustrated in this script, you can combine:

1. R code
2. R output
3. R graphics
4. text
5. hyperlinks
6. embedded graphics
7. headings, table of contents and other markup
8. mathematical formulas in LaTeX
9. and, with a bit of effort, interactive graphics and other widgets

in a document.

LaTeX is the most widely used language and environment for technical publishing.

In R Markdown, we have access to its language for mathematical formulas. For example, we can write:

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i \quad i = 1, \dots, n, \quad \epsilon \sim N(0, \sigma^2)$$

where $\beta_0 = 2$, $\beta_1 = 0.5$, $\sigma = 2$.

All you need in order to use R Markdown to produce HTML output is automatically installed when you install R Studio.

Installing R and R Studio

If you are running this script, you have probably already installed R and RStudio.

To install R visit this ‘mirror’ of the The Comprehensive R Archive Network and follow the instructions for your operating system: Mac OS X, Windows or Linux.

Read the information carefully. If you use Windows, install ‘Rtools’ as suggested. If you use Mac OS X, consider whether you need to install XQuartz.

After installing R, next install R Studio.

Once R Studio is installed, open it with the icon on your desktop. We will use the console in R Studio to install some packages.

Installing and loading packages

There are three main sources of packages for R:

1. Many packages come with R when you install it initially.
2. Most additional packages you might consider installing reside on ‘CRAN’, the Comprehensive R Archive Network, that provides relative stable versions of packages that have passed some automated tests, and
3. ‘github.com’ where most developers provide access to the latest beta versions of their packages or to packages that have not been sent for inclusion in CRAN.

The first package we install resides on CRAN and it is needed to install packages from GitHub.

Note that “#+ ” denotes a “chunk option” that modifies the behaviour of R until the end of the chunk. A chunk is ended by any line that starts with “#’ ” or with “#+ ”. The option “”’eval=FALSE” prevents R from running (evaluating) the chunk when you use Ctrl-Shift-K to produce an output file. We assume that all packages have already been installed interactively before using Ctrl-Shift-K.

If you haven’t installed the following packages yet, this is an opportunity to do so. You must run these lines ‘manually’ by using Ctrl-Enter with the cursor on the line.

```
# These lines will not be run when you render the script with Ctrl-Shift-K  
install.packages('devtools')
```

While we are at it we can install a few more packages from CRAN:

```
install.packages(c("car", "effects", "ggplot2", "Hmisc"))  
install.packages(c("knitr", "magrittr", "rgl", "rio", "rmarkdown", "readxl", "cv"))  
install.packages(c("latticeExtra"))  
install.packages('kableExtra')  
devtools::install_github('gmonette/spida2')  
devtools::install_github('gmonette/p3d')
```

Installing packages only needs to be done once every time you install a new version of R. You can update them occasionally, with:

```
update.packages()
```

Github packages need to be reinstalled periodically to get the latest updates:

```
# devtools::install_github('gmonette/spida2')  
# devtools::install_github('gmonette/p3d')
```

I expect these packages to be updated frequently during the course and you need to reinstall them to have access to the latest versions.

Each time you use R, you need to load the packages you need for that session with the ‘library’ command:

```
library(car)
```

Loading required package: carData

```
library(spida2)
library(lattice)
library(latticeExtra)
```

Reading data in a text file

The easiest formats to read are CSV files (comma-separated-value files which are easily created from Excel by saving a spreadsheet as a CSV file), and tab-delimited text files, usually with a ‘.txt’ extension.

You need to know the location of the file relative to the script file.

Be sure to set the working directory, which you can find with

```
getwd()
```

```
[1] "/home/georges/4939/www/files"
```

to the directory of the script file. In RStudio, use the menus:

Session > Set Working Directory > To Source File Location

Then you can refer to the file using a relative path. This is much better than an absolute path since you can send the script and the data set to a collaborator and they will be able to run the script after saving the script and the data set in a directory on their own computer.

To illustrate, we will download a data file and use for a brief exploration,

```
# You only need to download once, so this is in a code chunk that
# will not be run when you use Ctrl-Shift-K to render the file.
download.file("http://socserv.socsci.mcmaster.ca/jfox/Books/Applied-Regression-3E/datasets/Titanic.txt",
              "Titanic.txt")
```

This copies the file ‘Titanic.txt’ into the working directory. This file is a list of passengers on the Titanic and does not include the crew.

Note that all the data sets from the textbook can be downloaded this way. However, it is better to create a ‘data’ subdirectory and download them there.

```
list.files()
```

```
[1] "description.html"      "description.pdf"
[3] "description.R"         "Intro"
[5] "markdown_sample.pdf"  "markdown_sample.R"
[7] "markdown_sample.spin.R" "markdown_sample.spin.Rmd"
[9] "software_installation.html" "software_installation.R"
[11] "Titanic.txt"
```

You can also read a text file from the internet by using the URL but then you need to be connected to the internet whenever you read the file.

```
titanic <- read.table("Titanic.txt", header = T)
library(spida2)
library(lattice)
head(titanic) # first 6 lines
```

	survived	age	passengerClass
Allen, Miss Elisabeth Walton	yes	29.0000	1st

Allison, Miss Helen Loraine	no	2.0000	1st
Allison, Mr Hudson Joshua Creighton	no	30.0000	1st
Allison, Mrs Hudson J.C. (Bessie Waldo Daniels)	no	25.0000	1st
Allison, Master Hudson Trevor	yes	0.9167	1st
Anderson, Mr Harry	yes	47.0000	1st

	sex
Allen, Miss Elisabeth Walton	female
Allison, Miss Helen Loraine	female
Allison, Mr Hudson Joshua Creighton	male
Allison, Mrs Hudson J.C. (Bessie Waldo Daniels)	female
Allison, Master Hudson Trevor	male
Anderson, Mr Harry	male

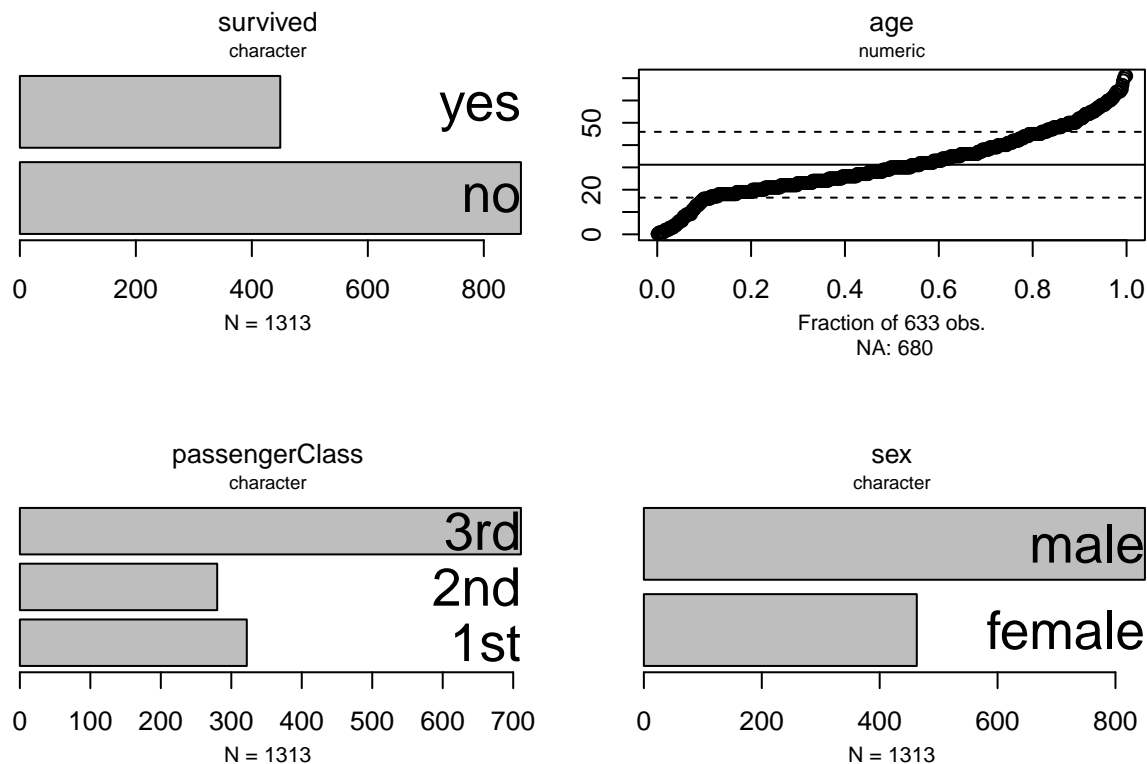
```
tail(titanic) # last 6 lines
```

	survived	age	passengerClass	sex
Zabour, Miss Tamini	no	NA	3rd	female
Zakarian, Mr Artun	no	NA	3rd	male
Zakarian, Mr Maprieder	no	NA	3rd	male
Zenn, Mr Philip	no	NA	3rd	male
Zievens, Rene	no	NA	3rd	female
Zimmerman, Leo	no	NA	3rd	male

```
dim(titanic) # rows and columns
```

```
[1] 1313 4
```

```
xqplot(titanic)
```



```
#
# frequency table
#
```

```
tab(titanic, ~ survived + sex + passengerClass)
```

```
, , passengerClass = 1st
```

	sex		
survived	female	male	Total
no	9	120	129
yes	134	59	193
Total	143	179	322

```
, , passengerClass = 2nd
```

	sex		
survived	female	male	Total
no	13	148	161
yes	94	25	119
Total	107	173	280

```
, , passengerClass = 3rd
```

	sex		
survived	female	male	Total
no	134	440	574
yes	79	58	137
Total	213	498	711

```
, , passengerClass = Total
```

	sex		
survived	female	male	Total
no	156	708	864
yes	307	142	449
Total	463	850	1313

```
#
# percentage within each 'sex by passengerClass' grouping
#
tab(titanic, ~ survived + sex + passengerClass, pct = c(2,3))
```

```
, , passengerClass = 1st
```

	sex		
survived	female	male	All
no	6.293706	67.039106	40.062112
yes	93.706294	32.960894	59.937888
Total	100.000000	100.000000	100.000000

```
, , passengerClass = 2nd
```

	sex		
survived	female	male	All
no	12.149533	85.549133	57.500000
yes	87.850467	14.450867	42.500000
Total	100.000000	100.000000	100.000000

```
, , passengerClass = 3rd
```

	sex		
survived	female	male	All
no	62.910798	88.353414	80.731364
yes	37.089202	11.646586	19.268636
Total	100.000000	100.000000	100.000000

```
, , passengerClass = All
```

	sex		
survived	female	male	All
no	33.693305	83.294118	65.803503
yes	66.306695	16.705882	34.196497
Total	100.000000	100.000000	100.000000

```
# nicer:  
tab(titanic, ~ survived + sex + passengerClass, pct = c(2,3)) %>%  
  round(1)
```

```
, , passengerClass = 1st
```

	sex		
survived	female	male	All
no	6.3	67.0	40.1
yes	93.7	33.0	59.9
Total	100.0	100.0	100.0

```
, , passengerClass = 2nd
```

	sex		
survived	female	male	All
no	12.1	85.5	57.5
yes	87.9	14.5	42.5
Total	100.0	100.0	100.0

```
, , passengerClass = 3rd
```

	sex		
survived	female	male	All
no	62.9	88.4	80.7
yes	37.1	11.6	19.3
Total	100.0	100.0	100.0

```
, , passengerClass = All
```

	sex		
survived	female	male	All
no	33.7	83.3	65.8
yes	66.3	16.7	34.2
Total	100.0	100.0	100.0

Note that the ‘%>%’ operator ‘pipes’ the output of the left-hand side (lhs) as the first argument of the function on the right.

In RStudio, you can type ‘%>%’ by pressing *Control-Shift-M*.

Visualizing frequencies

Barchart of frequencies

```
tab(titanic, ~ survived + sex + passengerClass)
```

```
, , passengerClass = 1st
```

	sex		
survived	female	male	Total
no	9	120	129
yes	134	59	193
Total	143	179	322

```
, , passengerClass = 2nd
```

	sex		
survived	female	male	Total
no	13	148	161
yes	94	25	119
Total	107	173	280

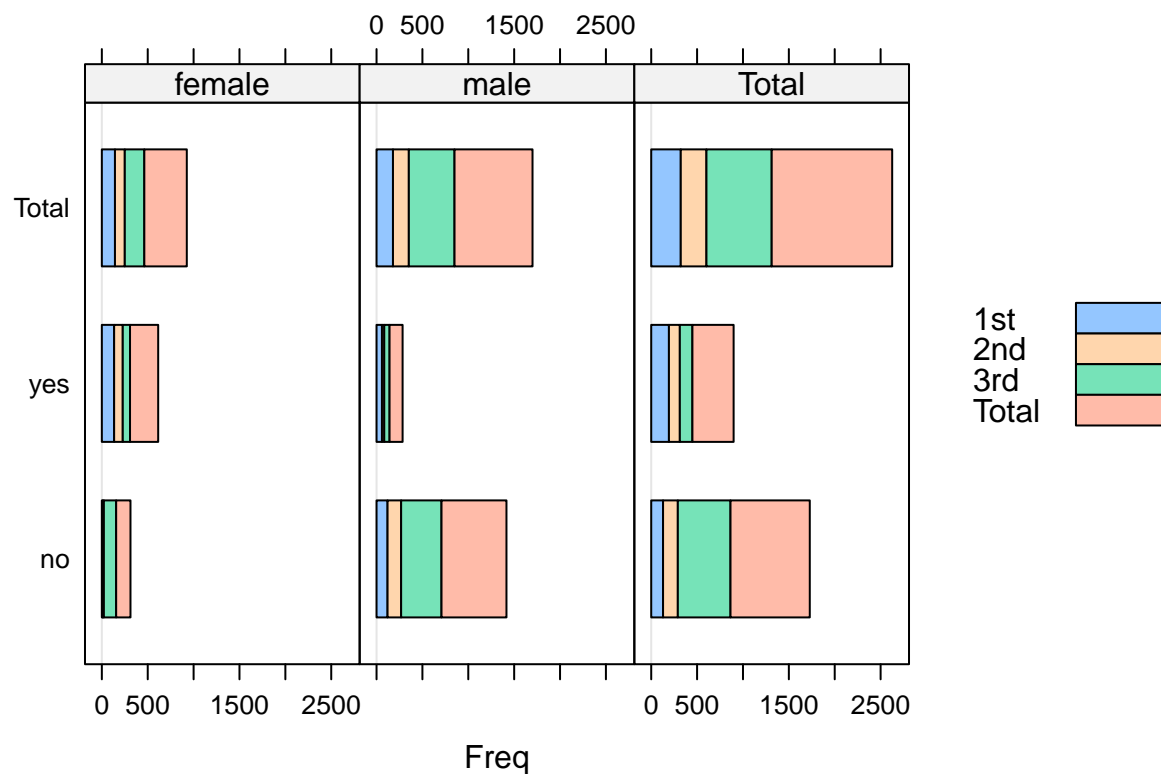
```
, , passengerClass = 3rd
```

	sex		
survived	female	male	Total
no	134	440	574
yes	79	58	137
Total	213	498	711

```
, , passengerClass = Total
```

	sex		
survived	female	male	Total
no	156	708	864
yes	307	142	449
Total	463	850	1313

```
tab(titanic, ~ survived + sex + passengerClass) %>%  
  barchart(auto.key=T)
```

Get rid of 'Total'

Using `tab_` instead of `tab` suppresses the “Total” margin which we don’t want to display since it is redundant in the graph. Using `tab__` suppresses both the “Total” margin the the “All” margin in conditional tables generated by using the arguments `pct` or `pr`.

```
tab_(titanic, ~ survived + sex + passengerClass)
```

```
, , passengerClass = 1st
```

```
sex
survived female male
no          9  120
yes       134   59
```


```
, , passengerClass = 2nd
```

```
sex
survived female male
no        13  148
yes       94   25
```


```
, , passengerClass = 3rd
```

```
sex
survived female male
no       134  440
yes       79   58
```

```
tab_(titanic, ~ survived + sex + passengerClass) %>%
  barchart(auto.key=T)
```



```
gd() # ggplot2-like appearance
tab_(titanic, ~ survived + sex + passengerClass) %>%
  barchart(
    auto.key = list(space = 'right', title='Class'),
    xlim = c(0, 800))
```



This is not really informative. We want to see the relative proportion of survivors in the different subgroups so we change the order of the variables. Also the labels 'yes' and 'no' are not informative themselves.

We change the order of the variables in the formula so the variable we want to see within bars comes last.

The first variable generates different bars within panels and the second variable generates the panels.

```
tab_(titanic, ~ sex + passengerClass + survived)
```

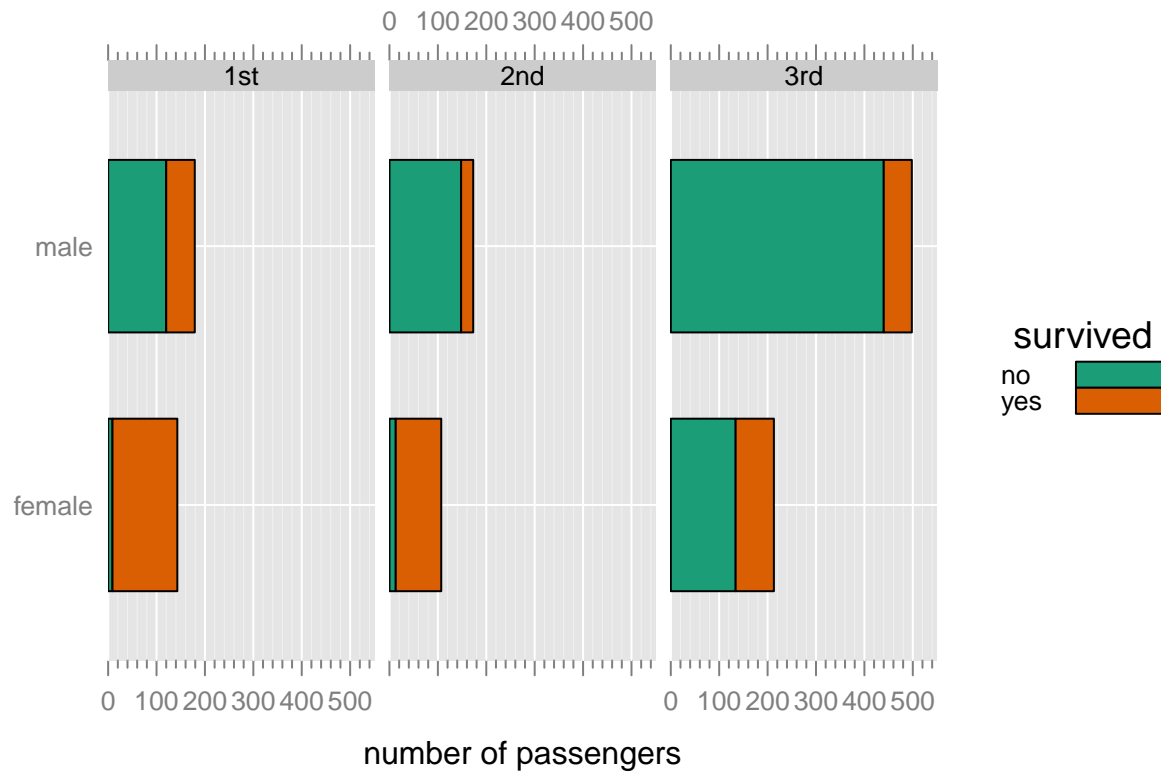
```
, , survived = no
```

```
      passengerClass
sex    1st 2nd 3rd
female  9  13 134
male   120 148 440
```

```
, , survived = yes
```

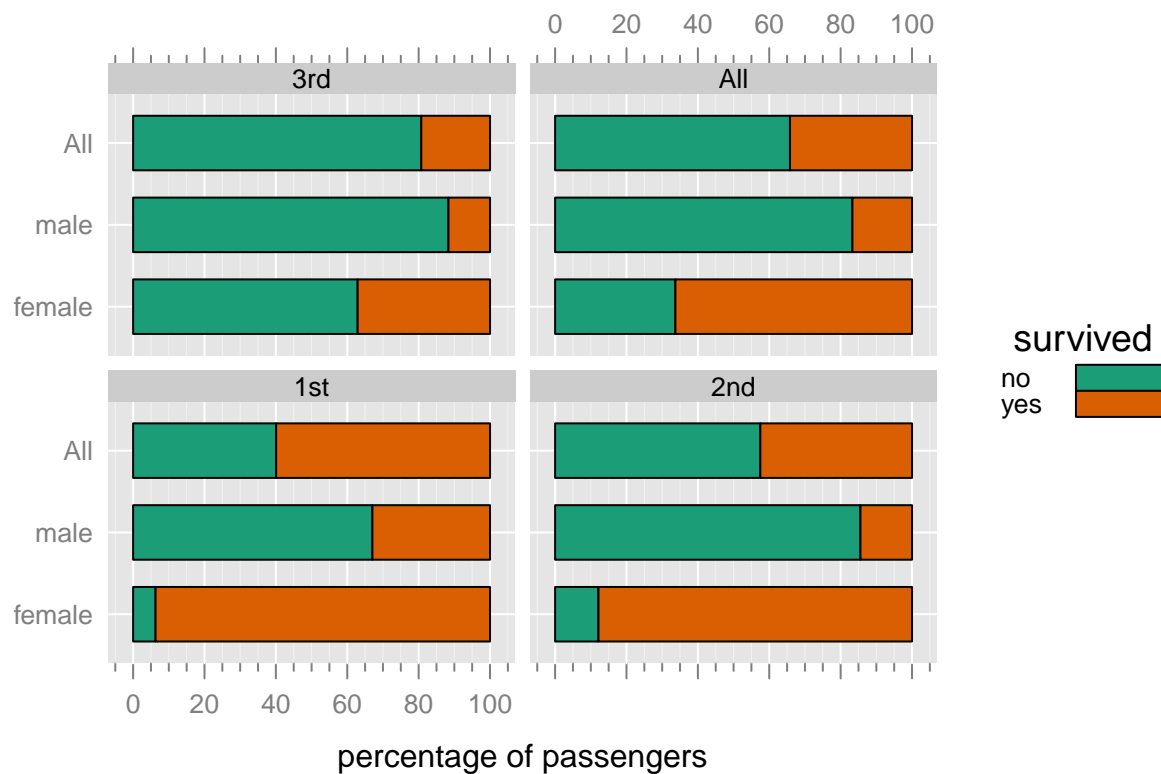
```
      passengerClass
sex    1st 2nd 3rd
female 134  94  79
male   59  25  58
```

```
tab_(titanic, ~ sex + passengerClass + survived) %>%
  barchart(xlab = 'number of passengers',
    xlim = c(0,550),
    auto.key=list(space='right',title='survived'))
```




But if we want to emphasize proportions we really want proportions

```
tab_(titanic, ~ sex + passengerClass + survived, pct = c(1,2)) %>%
  barchart(xlab = 'percentage of passengers',
    auto.key=list(space='right',title='survived'))
```



This shows overall proportions as well as within-gender proportions. To get rid of both 'All' and 'Total'. Use `tab__` instead of `tab`.

```
tab__(titanic, ~ sex + passengerClass + survived, pct = c(1,2)) %>%
  barchart(xlab = 'percentage of passengers', layout = c(1,3),
    auto.key=list(space='right',title='survived'))
```



experimenting:

```
tab__(titanic, ~ sex + passengerClass + survived, pct = c(1,2))
```


```
, , survived = no
```

```
      passengerClass
sex      1st      2nd      3rd
female 6.293706 12.149533 62.910798
male   67.039106 85.549133 88.353414
```

```
, , survived = yes
```


```
      passengerClass
sex      1st      2nd      3rd
female 93.706294 87.850467 37.089202
male   32.960894 14.450867 11.646586
```

```
tab__(titanic, ~ sex + passengerClass + survived, pct = c(1,2)) %>%
  barchart(ylab = 'percentage of passengers',
    horizontal = FALSE,
    ylim = c(0,100), layout = c(3,1),
    auto.key=list(space='right',title='survived'))
```



Looks wrong because order of colors are different in graph and in legend

```
tab__(titanic, ~ sex + passengerClass + survived, pct = c(1,2)) %>%
  barchart(ylab = 'percentage of passengers',
    horizontal = FALSE,
    ylim = c(0,100), layout = c(3,1),
    scales = list(x=list(rot=45)),
    auto.key=list(space='right',title='survived', reverse.rows = T))
```



more experimenting:

```
tab__(titanic, ~ passengerClass + sex + survived)
```


```
, , survived = no
```

```
sex
passengerClass female male
1st             9  120
2nd            13  148
3rd           134  440
```

```
, , survived = yes
```

```
sex
passengerClass female male
1st          134   59
2nd           94   25
3rd           79   58
```


```
gd()
tab__(titanic, ~ passengerClass + sex + survived) %>%
  barchart(ylab = 'number of passengers',
    ylim = c(0,500),
    horizontal = FALSE,
    auto.key=list(space='right',title='survived',
      reverse.rows=T))
```



Note that 'reverse.rows' get the key to show colors in the same order as in the plot.

```
tab__(titanic, ~ passengerClass + sex + survived) %>%
  barchart(ylab = 'number of passengers',
    ylim = c(0,500),
```

```
# box.width = rep(.9~c(2,1,1,1,1,1), each = 2),
  box.width = c(1,5,5,.1)/10,
  box.ratio = 2, stack = F,
  horizontal = FALSE,
  auto.key=list(space='right',title='survived', reverse.rows = T))
```



Questions

- What are the pros and cons of these various graphs?
- Can you find good visualizations of the Titanic data?
- What factors make a visualization effective?
- You've heard "a picture is worth a thousand words". But how about: "a picture is worthless unless it tells a good story." Do any of these plots tell a good story?
- What factors detract from the effectiveness of a visualization?
- Can you think of ways of presenting this data that would be more effective than the ones above?
- If you are familiar with 'ggplot2', use it to present this data?
- Try mosaic plots in the 'vcdExtra' package.