

# Mixed Models with R: Non-Linear Models

*Asymptotic Functions of Time*

Georges Monette

[random@yorku.ca](mailto:random@yorku.ca)

# Longitudinal models for IQ recovery after coma

Data: IQ tests on 200 post coma patients at QEH

Number of observations per patient:

Number of observations	1	2	3	4	5	Total
Number of patients	107	61	27	4	1	200

Method of analysis: Mixed models for longitudinal data  
Problem: Representing IQ recovery over time

## Some functions of time:

- linear
- quadratic
- higher polynomials
- splines
- exponential growth, decay
- exponential asymptotic growth
- periodic functions

Polynomials:

Linear function:

$$E(IQ) = \beta_0 + \beta_1 T$$

Quadratic:

$$E(IQ) = \beta_0 + \beta_1 T + \beta_2 T^2$$


Cubic

$$E(IQ) = \beta_0 + \beta_1 T + \beta_2 T^2 + \beta_3 T^3$$


Quartic:

$$E(IQ) = \beta_0 + \beta_1 T + \beta_2 T^2 + \beta_3 T^3 + \beta_4 T^4$$


Note the change in meaning of parameters as we move to higher order polynomials.  
Last parameter has global interpretation, previous parameters have interpretation at time = 0



5



6



7

## *Modeling individual trajectories*

A good strategy in longitudinal data analysis is to start with a plausible model for individual trajectories even if there is no data from any one individual to actually fit the model. If the data is unbalanced and you are willing to assume that the between-subject effect is close to the within-subject effect, then the estimation of individual trajectories 'borrows strength' from the between-subject model.

Within the limits imposed by sample size, we try to construct a model that:

1. captures the main theoretical properties of the phenomenon
2. preferably has interpretable parameters



To experiment with your model don't hesitate to simulate some plausible data with `locator` and play with model. Make an empty plotting surface, click to create some xy, the columns the names you want:


```
plot(0,0, xlim = c(0,800),  
      ylim = c(85,105), type = 'n')  
iqsim.ex <- locator( 10 , type = 'p')  
iqsim.ex  
iqsim.ex <- as.data.frame( iqsim.ex )  
iqsim.ex  
names( iqsim.ex ) <- c('days','iq')  
iqsim.ex <- iqsim.ex[  
  order(iqsim.ex$days),]
```

However, we'll use precooked data:

```
> data( iqsim )      # from spida
> iqsim
```

	days	iq
1	30.9375	89.07734
2	73.1250	91.74573
4	101.2500	94.12407
3	104.3750	91.28166
5	198.1250	96.73445
6	249.6875	96.03835
7	249.6875	96.44441
8	285.6250	97.89462
9	323.1250	97.72059
10	335.6250	98.47470

```
plot( iq ~ days ,  
      iqsim, pch = 16, xlim = c(0,800),  
      ylim = c(85,105) )
```



# Fitting a line:

```
> fit.lin <- lm ( iq ~ days, iqsim )
> summary( fit.lin )
. . . . .
```

## Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	89.540151	0.759021	117.97	2.98e-1
days	0.027739	0.003429	8.09	4.03e-0

Residual standard error: 1.133 on 8 degrees of f  
Multiple R-squared: 0.8911, Adjusted R-squar  
F-statistic: 65.45 on 1 and 8 DF, p-value: 4.02


## Graphing a fitted line

We would like to show the predicted value over the whole range of days in the graph, not just the values that were observed. For a straight line we could just use `abline`. With curved lines we need a different approach. So we create a prediction data frame with the one predictor variable.

```
> pred <- expand.grid( days = seq( -20, 850, by = 100 ),
> pred$iq.lin <- predict( fit.lin, pred )
> some( pred )
```

	days	iq.lin
94	73	91.56510
137	116	92.75788
203	182	94.58865
.	.	.
746	725	109.65094
753	732	109.84511

```
> lines( iq.lin ~ days, pred,
          col = 'blue', lwd = 2)
```



14

Doesn't make much sense!

Let's try a quadratic

```
> fit.quad <- lm( iq ~ days + I(days ^2),  
> summary( fit.quad )
```

```
. . . .
```

Coefficients:


	Estimate	Std. Error	t value	Pr
(Intercept)	8.780e+01	1.170e+00	75.020	1.
days	5.506e-02	1.535e-02	3.586	
I(days^2)	-7.324e-05	4.036e-05	-1.815	

Residual standard error: 0.9988 on 7 degrees of

Multiple R-squared: 0.9259, Adjusted R-squar

F-statistic: 43.75 on 2 and 7 DF, p-value: 0.00


```
> pred$iq.quad <- predict( fit.quad, pred )  
> lines( iq.quad ~ days , pred, col = 'red',  
>
```



16



With a quadratic, what goes up must come down ... the went up! Maybe a cubic makes more sense:



Exasperated we decided to go all the way with a polynomial degree 8:

Thanks to John, this option


```
> p8 <- function( x ) poly( x, 8, raw = TRUE)
> fit.high <- lm( iq ~ p8( days ), iqsim )
> summary(fit.high)    # look at R-Squared!!
```

. . . .  
Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	1.314e+03	1.665e+02	7.889	0.0803	.
p8(days)1	-9.135e+01	1.238e+01	-7.381	0.0857	.
p8(days)2	2.548e+00	3.436e-01	7.414	0.0854	.
p8(days)3	-3.596e-02	4.829e-03	-7.447	0.0850	.
p8(days)4	2.878e-04	3.846e-05	7.482	0.0846	.
p8(days)5	-1.362e-06	1.811e-07	-7.518	0.0842	.
p8(days)6	3.777e-09	4.999e-10	7.555	0.0838	.
p8(days)7	-5.674e-12	7.476e-13	-7.590	0.0834	.
p8(days)8	3.567e-15	4.679e-16	7.624	0.0830	.

---  
Residual standard error: 0.2871 on 1 degrees of freedom  
Multiple R-squared: 0.9991, Adjusted R-squared: 0.9921  
F-statistic: 142.8 on 8 and 1 DF, p-value: 0.06463

*An almost perfect fit!*



19


- A perfect fit to the data
- But a very poor fit to the 'population'
- An example of overfitting and loss of validity


*The remedy:*

Use a model that captures characteristics of the process  
Don't just use a high order polynomial to get a good em


Presumably, under typical circumstances, recovery reaches a peak  
after a while. We need a model that rises at first and then comes  
out.


*Exponential growth or decay*






*Exponential decay*











*Exponential asymptotic growth*

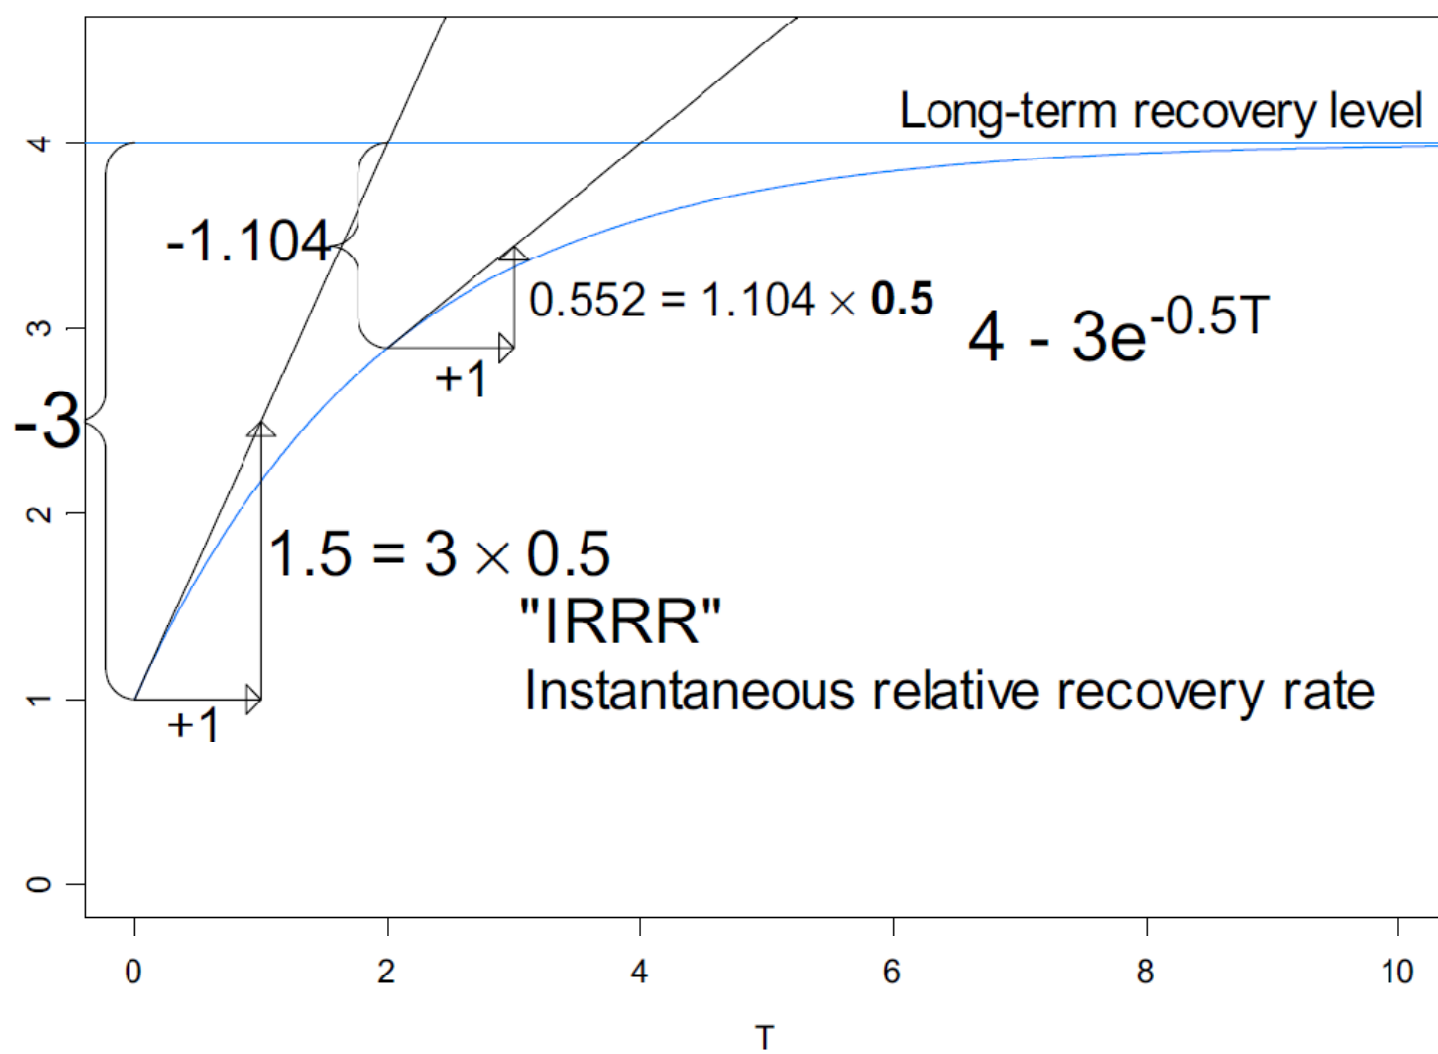




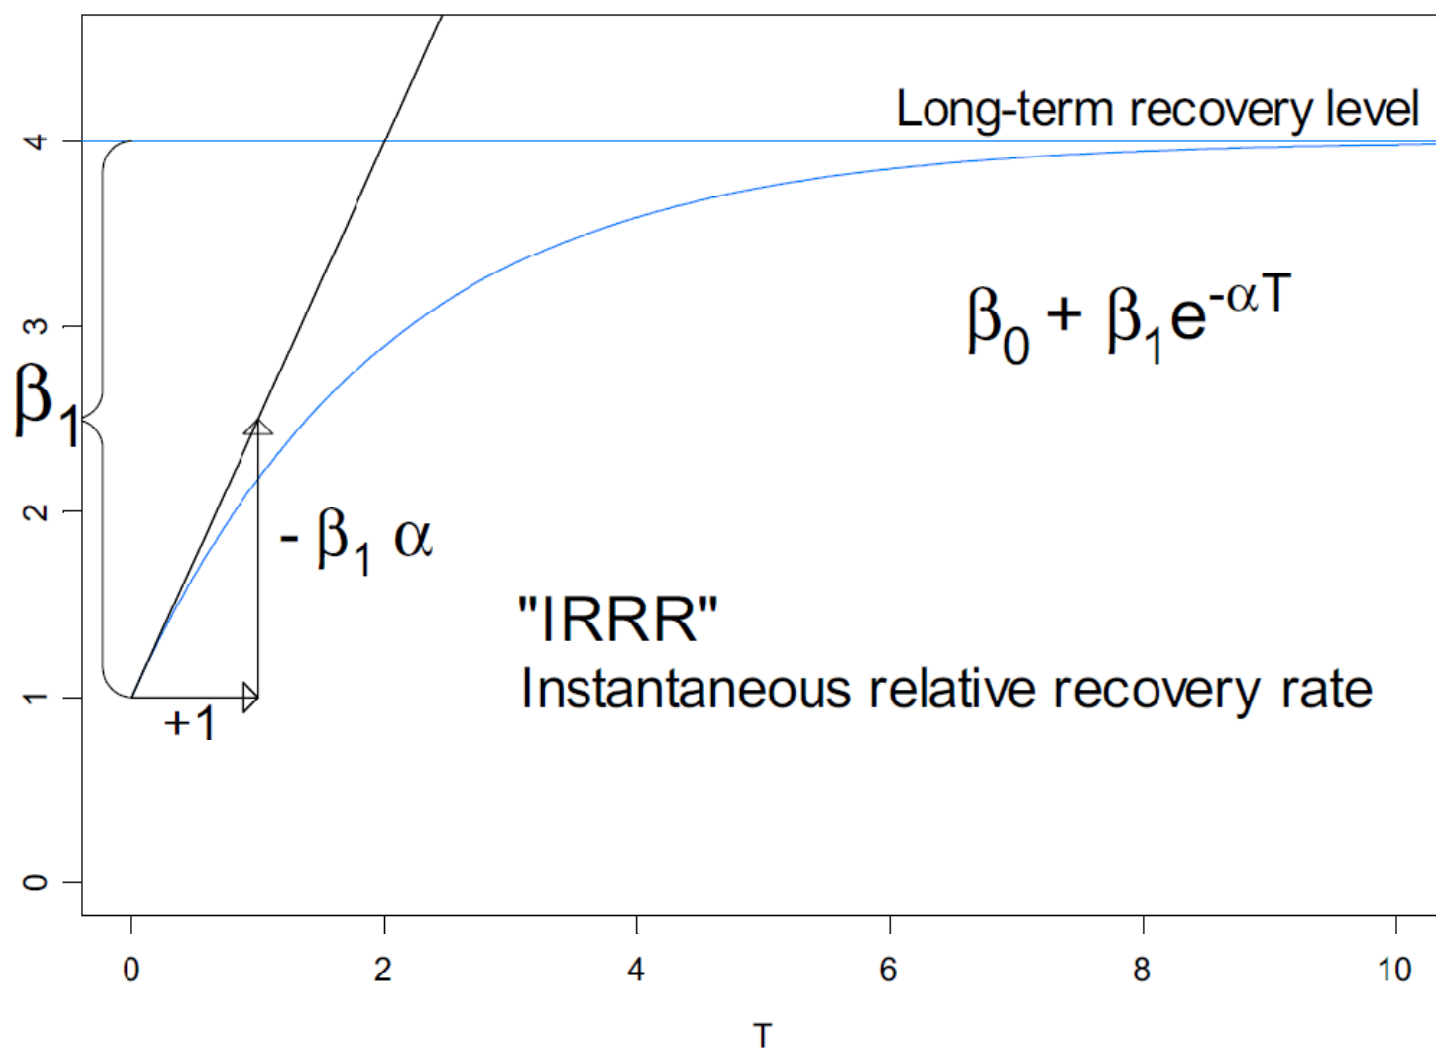


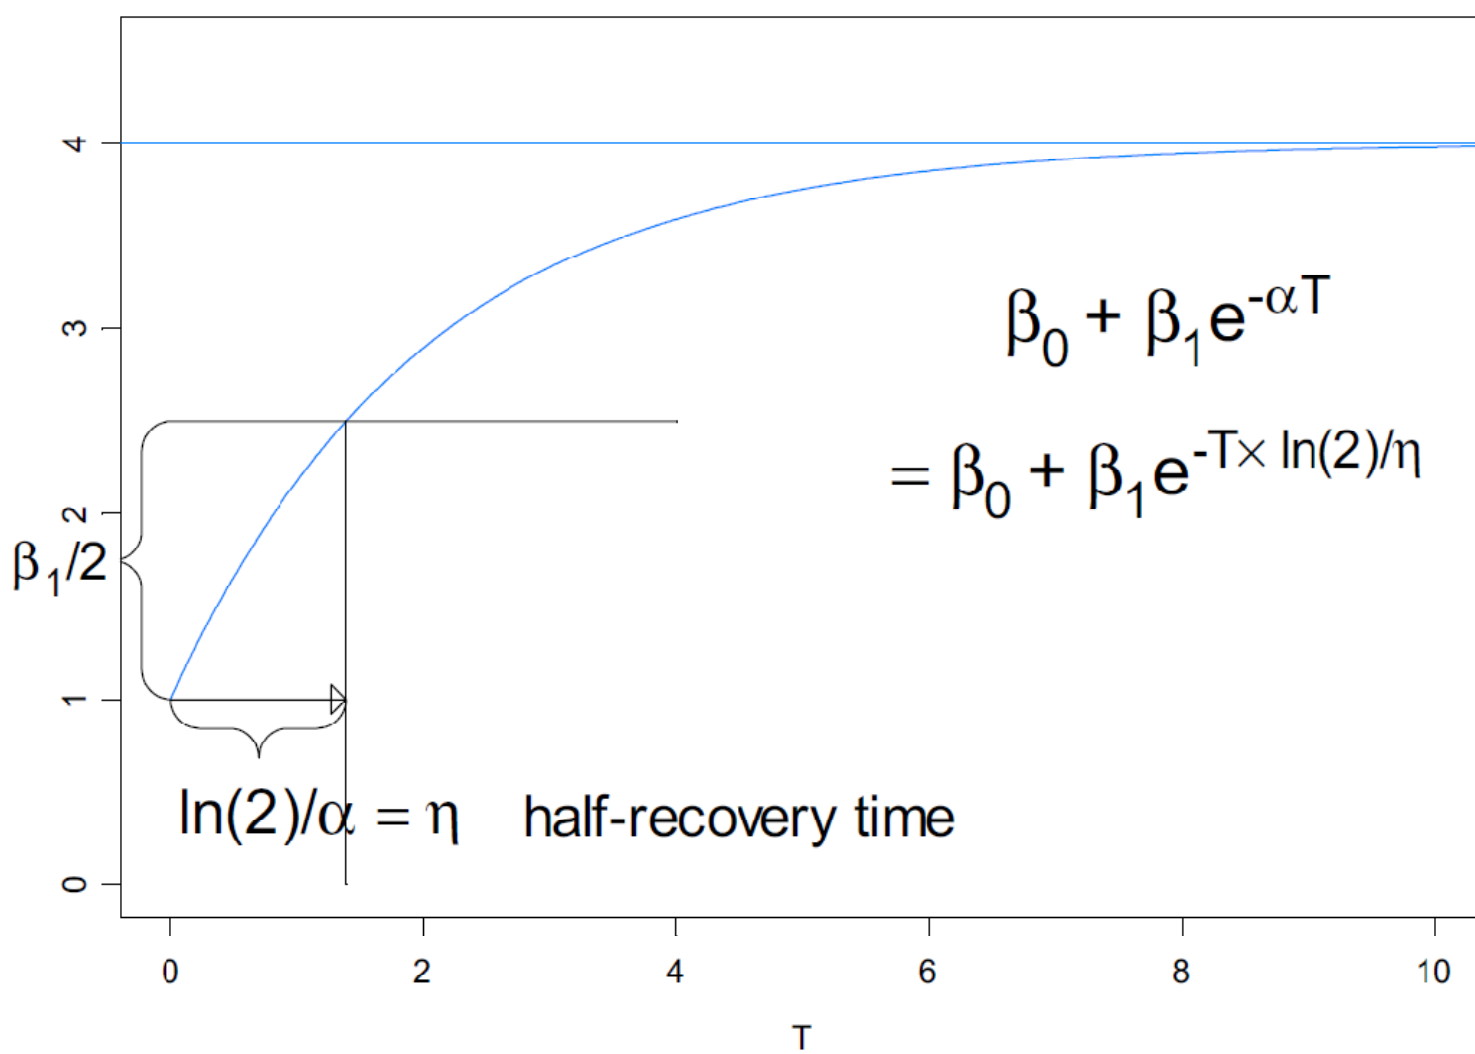






## Using half-recovery time instead (half-life)







## *Fitting a non-linear growth curve model:*

$$Y_t = \beta_0 + \beta_1 \exp(-\alpha T_i) + \varepsilon_i \quad \varepsilon_i \text{ independent } N(0, \sigma)$$

Later we will use nlme for longitudinal data with more than one subject. With just one subject we use 'nls', which is to 'nls' 'lm' is to 'lme'.

The syntax for fitting a non-linear model is very similar to fitting a linear model with three differences.

1. With a linear model we only need to specify the *model formula*. We don't need to say anything about the *parameters* because it is understood that there is exactly one parameter for each regressor (some predictors will have more than one parameter) and each parameter multiplies its regressor. The *model formula* for a non-linear model needs to specify the *parameters*.

the parameters and the regressor.

2. The algorithm for fitting is iterative and needs starting values which you generally need to supply.
3. In non-linear mixed effects models – with **nlme** – the non-linear model are themselves be modelled as linear models potentially based on other predictors. This allows the non-linear model to be simpler since it only has to capture the essentially non-linear aspects of the data. Another advantage is that this formulation is easy to implement numerically, i.e. it's less work for the computer.

# Growth curve model:

$$Y_i = \beta_0 + \beta_1 \exp(-\alpha T_i) + \varepsilon_i \quad \varepsilon_i \text{ i.i.d. } N(0, \sigma^2)$$

Non-linear model formula:

```
iq ~ b0 + b1*exp(-alpha*days)
```

The formula contains references to data: **iq**, **days** that are found in the **iq** data frame.

Parameters: **b0**, **b1**, **alpha** that need starting values.

Finding starting values: best way: sketch and understand the data and infer plausible parameters.

From graphs I would guess:

```
list( b0 = 100, b1 = -20, alpha = 0
```

How did we get these values?

**b0** is the long-run level, **b1** is the relative deficit at  
**alpha** is the daily proportion of lost iq recovered. I  
might take 100 days for a half recovery of 0.5, so di  
suggests roughly 0.005 per day.

Call in R:

```
nls( iq ~ b0 + b1*exp(-alpha*days),  
      start = list( b0 = 100, b1 = -20  
                    alpha = 0.005 ) )
```

## *R code and output:*

```
> fit.nl <- nls ( iq ~ b0 + b1*exp( -a*days  
+               start= list(b0 = 100, b1 = -30, a  
> summary( fit.nl )
```

Formula:  $iq \sim b0 + b1 * \exp(-a * \text{days})$

Parameters:

	Estimate	Std. Error	t value	Pr(> t )	
b0	99.906891	2.393470	41.741	1.18e-09	**
b1	-12.847352	1.620520	-7.928	9.66e-05	**
a	0.005820	0.002956	1.969	0.0897	.

---


Residual standard error: 0.9738 on 7 degrees of

Number of iterations to convergence: 4


Achieved convergence tolerance: 7.917e-06

```
> pred$iq.nl <- predict( fit.nl, pred )
> plot( iq ~ days , iqsim, pch = 16,
+       xlim = c(0,800), ylim = c(85,105) )
> lines( iq.nl ~ days , pred, col = 'black', lwd
> coef( fit.nl )
               b0               b1               a
99.906891397 -12.847351723    0.005819758
> abline( h = coef(fit.nl)[1], col = 'gray',
```

Asymptotic growth curve:



Polynomials:





## *Alternative: Transforming Time*

What difference does it make if we turn our non-linear model into a linear model by transforming time:

$$ttime = \exp\{-0.0056 \times \text{days}\}$$

As  $\text{days} \rightarrow \infty$ ,  $ttime(\text{days}) \rightarrow 0$ , as  $\text{days} \rightarrow 0$ ,  $ttime(\text{days}) \rightarrow \infty$

```
> ttime <- function( x ) exp( -0.0056 * x )
> fit.lin <- lm( iq ~ ttime( days ) )
> summary(fit.lin)
```

**Coefficients:**

	Estimate	Std. Error	t value
(Intercept)	99.9224	0.5628	177.54
ttime(days)	-12.8535	1.2508	-10.28

Residual standard error: 0.9109 on 8 degrees of freedom  
Multiple R-squared: 0.9296, Adjusted R-squared: 0.9214  
F-statistic: 105.6 on 1 and 8 DF, p-value: 6.92e-06

*Compare coefficients from tranformed fit and from non-linear fit*

Coefficients: (transformed)

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	99.9224	0.5628	177.54	1.13e-15
ttime(days)	-12.8535	1.2508	-10.28	6.92e-06

Parameters: (non-linear)

	Estimate	Std. Error	t value	Pr(> t )
b0	99.906891	2.393470	41.741	1.18e-09 ***
b1	-12.847352	1.620520	-7.928	9.66e-05 ***
a	0.005820	0.002956	1.969	0.0897 .

- The **estimated parameters** are almost the same but the linear fit using transformation reports **much smaller SEs** than the non-linear fit.
- Why? The linear fit is not taking into account the uncertainty stemming from the unknown asymptote.
- Note that the biggest difference in SE occurs for the asymptote. Unless you know the asymptote – where the curve gets very flat – the estimate of the asymptote is very sensitive to the estimate of curvature.
- When reviewing work that used transformations consider whether a non-linear fit would have been more honest. Note that the transformation is free if it is an intented transformation on the scale: e.g. log(Salary).


# Summary:

We fit an asymptotic non-linear growth curve to PIQ as a function of time post Coma.


We then try the same model with VIQ in the hope of comparing the two models, but the VIQ model does not converge. We explore various remedies of non convergence and eventually decide on a minor reparametrization. It works! We then use the same model on PIQ and compare the two models.

These two models are 'univariate' multivariate models, looking at one response at a time. To get p-values in the comparison of the two responses, we need to do more. One possibility is bootstrapping, which we don't explore. The other is to exploit multilevel (with 3 levels) modeling in nlme to fit something close to (but not exactly) a multivariate model. This is done at the end of the Lab script.


# Recovery of post-coma IQ



First 3 years:



45



46

## *Using nlme*

The nlme model is specified like a hierarchical model, and you can mix variable levels.

Example:

```
nlme( piq ~ b0 + b1*exp(-a*dayspc),  
      data = iq,  
      fixed = list( b0 ~ 1+sqrt( dcoma  
        b1 ~ 1,  
        a ~ 1),  
      random = list( id = b0 ~ 1),  
      start = list( fixed = c(  
        100, 0, -20., .05)),  
      control = list( maxIter = 100,  
        returnObject = TRUE),  
      verbose = T)
```

*The code one line at a time:*

```
piq ~ b0 + b1*exp(-a*dayspc)
```

A *non-linear model formula* with *regressors* and parameters. In this example, it's the Level 1 model. In general, you can add Level 2 *regressors* in this model. If you want to use a categorical variable, you need to use it through its dummies: e.g.

```
b.sex * (sex=="Female")
```

where **b.sex** is the parameter multiplying the indicator variable for Female.

```
data = iq,
```

data frame as usual



```
fixed = list(  
  b0 ~ 1+sqrt( dcoma ),  
  b1 ~ 1+sqrt( dcoma ),,  
  a ~ 1)
```

A list of *linear model formulas*, one for each parameter. In the model, the parameter **a** is assumed to have the same value across the population, **b0** and **b1** are assumed to depend through the model on **sqrt(dcoma)**. This transformation incorporates the assumption that an extra day of coma after, say, 3 days has a greater impact than an extra day after 50 days. The square root transformation was chosen by examining visual plots of the data, which was somewhat arbitrary. Also it is an oversimplification to assume that **a** is a constant across the population.

```
random = list( id = list( b0 ~ 1, b1
```

Specify the parameters that are assumed to vary random  
id to id. Note that **b0** is the asymptotic level but it is  
constant added to all observations.

```
start = list( fixed =  
  c( 100, -10, -20., -1, .05) ) )
```

This is the challenging part that rewards a good understanding of the parameters of the model. Recall the fixed portion above:

```
fixed = list(  
  b0 ~ 1+sqrt( dcoma ),  
  b1 ~ 1+sqrt( dcoma ),  
  a ~ 1)
```

The starting values are listed in the same order as the parameters of the 'fixed' portion of the model. Generally, it is good to have plausible starting values. Draw a sketch and make your own guesses. Here, our starting model is:

```
b0 = 100 - 10 * sqrt(dcoma)  
b1 = -20 - 1 * sqrt(dcoma)  
a = 0.05
```

```
control = list( maxIter = 100,  
               returnObject = TRUE)
```

Increases the default number of iterations from 50 to 100. It also returns the last fit even if there is no convergence. Watch the console to use this shortly.

```
verbose = T
```

This shows information on each **PNLS** and **LME** step. Press Ctrl-W in the R console to get unbuffered output and watch a frequently exciting show.

## *Fitting the model:*

```
> fit.nlme <- nlme(  
+   piq ~ b0 + b1*exp(-a*dayspc),  
+   data = iq,  
+   fixed = list(  
+     b0 ~ 1 + sqrt(dcoma) ,  
+     b1 ~ 1 + sqrt(dcoma) ,  
+     a ~ 1),  
+   random = list( id = list( b0 ~ 1, b1~ 1  
+   control = list( maxIter = 200, returnObje  
+   start = list(  
+     fixed = c(100, -10, -10, 0,.05)),  
+   control = list( maxIter = 100, returnObje  
+   verbose = TRUE)
```

. . . . [Omitting Output on Iterations 1 to 3]

**\*\*Iteration 4**

LME step: Loglik: -1287.679 , nlm iterations: 1

reStruct parameters:

id1	id2	id3
1.515913	0.949715	23.921793

PNLS step: RSS = 15018.94

fixed effects:97.0948 -1.24521 -11.1453 -3.24829

iterations: 7

This was pretty quick co

Convergence:

fixed	reStruct
1.312661e-06	7.021607e-04

> summary( fit.nlme )

Nonlinear mixed-effects model fit by maximum lik

Model: piq ~ b0 + b1 \* exp(-a \* dayspc)

Data: iq

AIC	BIC	logLik
2593.358	2627.577	-1287.679

Random effects:  
Formula: list(b0 ~ 1, b1 ~ 1)  
Level: id  
Structure: General positive-definite, Log-Cholesky parametrization

	StdDev	Corr
b0.(Intercept)	13.769293	b0.(I)
b1.(Intercept)	2.605835	-0.994
Residual	6.736055	

worries me a bit  
I might try to reparametr

Fixed effects: list(b0 ~ 1 + sqrt(dcoma), b1 ~ 1 + sqrt(dcoma), a ~ 1)

	Value	Std.Error	DF	t-value
b0.(Intercept)	97.09476	2.036582	127	47.67536
b0.sqrt(dcoma)	-1.24521	0.480486	127	-2.59157
b1.(Intercept)	-11.14530	3.208072	127	-3.47414
b1.sqrt(dcoma)	-3.24829	1.076749	127	-3.01676
a	0.00825	0.001651	127	4.99579

Correlation:

	b0.(I)	b0.s()	b1.(I)	b1.s()
b0.sqrt(dcoma)	-0.724			
b1.(Intercept)	-0.596	0.463		
b1.sqrt(dcoma)	0.463	-0.455	-0.789	
a	-0.309	0.013	0.092	-0.380

Standardized Within-Group Residuals:

Min	Q1	Med	Q3
-3.332408193	-0.365688335	0.009002275	0.382738703

2.30311

Number of Observations: 331  
Number of Groups: 200

> *An interesting calculation:*

Between subject SD of 'true' IQ	13.769
Within subject between test SD of IQ	6.736
Population SD of IQ	$\sqrt{13.769^2 + 6.736^2} =$
Test-retest reliability of IQ $= \frac{\text{Variance in True Score}}{\text{Variance of Observed Score}}$	$\frac{13.769^2}{15.328^2} = 0.8$



## *How does fitting work?*

See Pinheiro and Bates (2000) and Lindstrom and Bates (1980) for a description. It's a clever blend of available tools. Bates and Watts (1988) *Non-linear regression analysis and its applications* deals with non-linear models for independent data which is adapted to situation where the variance-covariance is known. So we have tools for non-linear models when the variance-covariance is known. And we have tools for linear mixed models (Lindstrom and Bates 1980) which give estimates of parameters in a non-linear model by fitting an approximating linear model.

The algorithm keeps repeating 2 steps until convergence.

1) PNLIS step: Given an estimate of  $G$  and  $R$ , estimate fixed parameters and random effects using a *penalized non-linear least squares* algorithm.

2) LME step: Given estimates of fixed parameters and random effects, construct an approximating linear model and estimate it in R with lme.


Keep repeating (1) and (2) until the estimates don't change.

# Some diagnostics for PIQ

```
> plot( fit.nlmme, resid(. , type = 'p') ~ fi
+       id = .05)
```



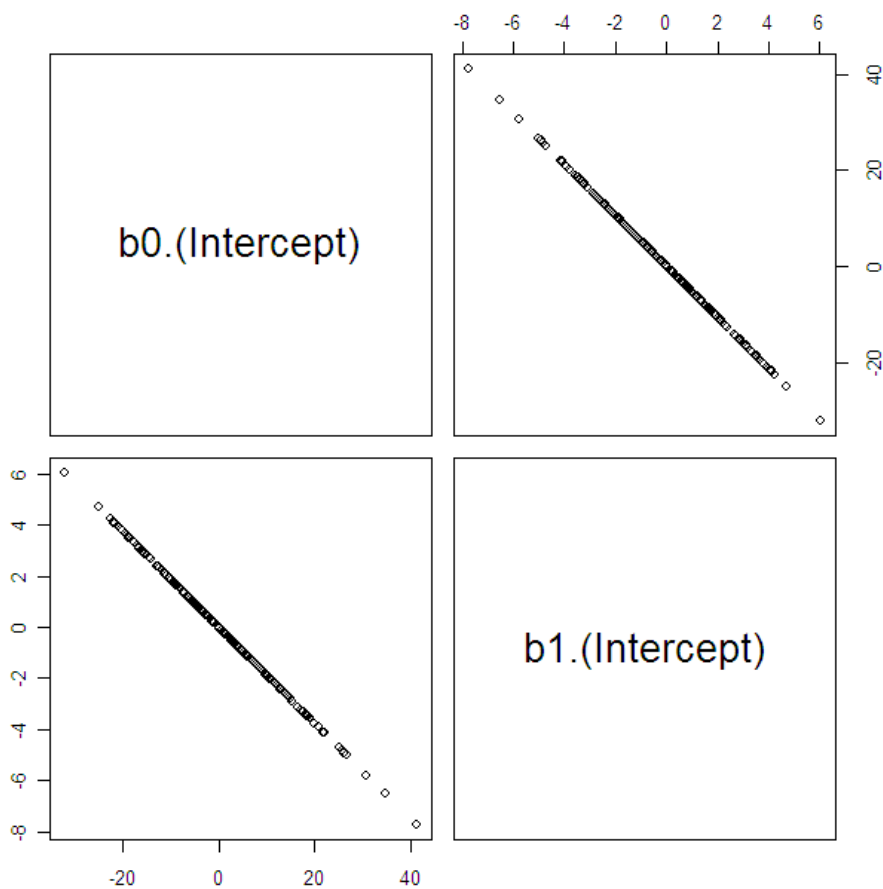
```
> plot( fit.nlme, sqrt( abs( resid(. ,type='
+ ~fitted(.), id = .05)
```



```
> plot( ranef( fit.nlme )) # output omitted
```

60

```
> pairs( ranef( fit.nlme ))
```



This plot shows the near singularity of the G matrix showing the strong correlation between **b0** and **b1**. It suggests that the recovery level, possibly related to pre-trauma intelligence, confer so large a benefit in the early stages of recovery.

## *Fitting VIQ*

```
> fit.nlme.viq <- nlme(  
+   viq ~ b0 + b1*exp(-a*dayspc),  
+   data = iq,  
+   fixed = list( b0 ~ 1 + sqrt(dcom  
+                 b1 ~ 1 + sqrt(dcom  
+                 a ~ 1)),  
+   random = list( id =  
+                 list( b0 ~ 1, b1~ 1 )),  
+   start = list(  
+     fixed = c(100, -.3, -10, -5,  
+     control = list( maxIter = 100,  
+       returnObject = T),  
+     verbose = T)
```

## **\*\*Iteration 1**

**LME step:** Loglik: -1246.109 , nlm iterations: 20

reStruct parameters:

id1	id2	id3
-0.31704425	-0.03919651	1.58004210

**PNLS step:** RSS = 8190.372

fixed effects: 98.3784 -0.398053 -11.1171 -3.81443

iterations: 7

### **Convergence:**

fixed	reStruct
0.6045864	0.3340419

We hope these numbers  
will get very small

Then it repeats:

## **\*\*Iteration 2**

**LME step:** Loglik: -1242.047 , nlm iterations: 10

reStruct parameters:

id1	id2	id3
-0.6005839	-1.1484260	0.3314853

**PNLS step:** RSS = 29302.31

```
fixed effects:93.3904  -0.500735  -2.79404  -5.93326
iterations: 7
```

Convergence:

```
fixed reStruct
2.978855 2.355074
```

Much, much later:

**\*\*Iteration 97**

LME step: Loglik: -1223.248 , nlm iterations: 11

reStruct parameters:

```
id1 id2 id3
-0.3076645 -0.9642632 0.4746019
```

PNLS step: RSS = 13520.03

fixed effects:96

iterations: 7

Convergence:

```
fixed reStruct
0.6185455 1.3370712
```



```
**Iteration 98
LME step: Loglik: -1239.934 , nlm iterations: 11
reStruct parameters:
      id1      id2      id3
-0.6920771 -0.9315449  0.3989076

PNLS step: RSS =  9035.715
fixed effects:98.3663  -0.345367  -8.12821  -2.46351
iterations: 7

Convergence:
      fixed reStruct
1.619486 0.572147

**Iteration 99
LME step: Loglik: -1223.284 , nlm iterations: 11
reStruct parameters:
      id1      id2      id3
-0.3084674 -0.9647008  0.4741239

PNLS step: RSS =  13512.85
fixed effects:96.0224  -0.381974  -9.02973  -5.45431
iterations: 7
```

These are the esti  
effects at iteration

```
Convergence:
      fixed reStruct
0.6181109 1.3332673
```

These numbers are just  
bouncing around and not  
getting smaller

```
**Iteration 100
LME step: Loglik: -1239.936 , nlm iterations: 11
reStruct parameters:
      id1      id2      id3
-0.6921102 -0.9316034  0.3988804

PNLS step: RSS =  9033.863
fixed effects:98.3667  -0.345443  -8.13618  -2.46378
iterations: 7
```

Have a close look  
number have been  
past few iterations

```
Convergence:
      fixed reStruct
1.6167060 0.5714619
Warning message:
In nlme.formula(viq ~ b0 + b1 * exp(-a * dayspc), data
= list(b0 ~  :
Maximum number of iterations reached without converg
```

Our model didn't converge.

## *What to do?*

Looking carefully at the output we notice that:

- 1) The convergence criteria are not slowly getting smaller, bouncing around, and
- 2) Some of the fixed parameters are stuck in an oscillation.

Possible actions:

- 1) Increase maxIter and come back much later. In this case, it might help,
- 2) Simplify the model, particularly the RE model. Here, it might work and it would work – but it would require making assumptions we would prefer to avoid for now.
- 3) Consider whether the model is well identified. Are there parameters whose estimates are far from expected? Are they constantly changing in a given direction as iterations progress?

- 3) When the process is stuck in a cycle, as it is here, with parameters oscillate and how do they oscillated together visualize what this implies about the fitted surface. It mean that the leverage and residuals of some points change from iteration to iteration. Such points pull the fit towards themselves in one iteration but in the new fit they lose influence and let go of the fitted surface for the next iteration. This doesn't happen with OLS because leverage does not depend on the previous fit so leverage can change from one fit to the next. Points involved tend to be outliers.

Finding problematic points:

- 1) Use `maxIter` to stop the iteration at each point in the 2-cycle, use an odd iteration and an even iteration. Save the object in each case. Plot the residuals of one fit against the other.

residuals of the other. The points with very different residuals are likely to be the culprits. You can also compare the residuals of each fit attempts to approximate the data.

2) Use your knowledge of the data to isolate some known good data.

We note that the distribution of `dcoma` is highly skewed. We try dropping anyone with `dcoma > 100`. Note `dcoma` is a subject variable and extreme values are likely to have high influence. With the asymptotic model for `dayspc`, influence may be high with small values being very influential and large values being less so. After further experimenting we also see that `b1`, the deficit at `dayspc = 0` is bound with the estimate of `a`. A large positive `a` produces a very negative value of `b1`. This leads us to the realization that attempting to estimate deficit immediate arousal is not feasible since there is little IQ data that easily move the origin to say one month after arousal by

changing `dayspc` to `dayspc - 30` in the non-linear  
this works we can do the same for PIQ to have compar

```
> fit.nlme.viq2 <-nlme( viq ~ b0 + b1*exp(-a*(da
+       data = iq,
+       fixed = list( b0 ~ 1 + sqrt(dcoma) ,
+                     b1 ~ 1 + sqrt(dcoma) ,
+                     a ~ 1),
+       random = list( id = list( b0 ~ 1, b1~
+       start = list(
+       fixed = c(100, -.3, -10, 0,.3)),
+       control = list( maxIter = 100, returnO
+       verbose = T,
+       subset = dcoma < 100)
```

which converges in 3 iterations:

• • • • •

```

**Iteration 3
LME step: Loglik: -1225.668 , nlm iterations
reStruct parameters:
      id1      id2      id3
-0.8286352 12.2429072 59.7420069

PNLS step: RSS = 9725.302
  fixed effects:99.2084 -0.561859 -6.79895
0.0214789
iterations: 7

Convergence:
      fixed      reStruct
1.054542e-07 1.588532e-01
>
> summary( fit.nlme.viq2 )
Nonlinear mixed-effects model fit by maximum
likelihood

```

Model: viq ~ b0 + b1 \* exp(-a \* (dayspc -  
Data: iq  
Subset: dcoma < 100  
AIC BIC logLik  
2469.336 2503.363 -1225.668

Random effects:  
Formula: list(b0 ~ 1, b1 ~ 1)  
Level: id  
Structure: General positive-definite, Log-Cholesky  
parametrization

	StdDev	Corr
b0.(Intercept)	1.254731e+01	b0.(I)
b1.(Intercept)	2.640292e-05	0
Residual	5.478719e+00	



Fixed effects: list(b0 ~ 1 + sqrt(dcoma), b1  
sqrt(dcoma), a ~ 1)

	Value	Std.Error	DF	t-val
b0.(Intercept)	99.20845	1.7262297	196	57.471
b0.sqrt(dcoma)	-0.56186	0.4940133	123	-1.137
b1.(Intercept)	-6.79895	2.2442154	123	-3.029
b1.sqrt(dcoma)	-1.87309	0.7804948	123	-2.399
a	0.02148	0.0044938	123	4.779

Correlation:

	b0.(I)	b0.s()	b1.(I)	b1.s()
b0.sqrt(dcoma)	-0.777			
b1.(Intercept)	-0.418	0.332		
b1.sqrt(dcoma)	0.312	-0.343	-0.860	
a	-0.148	-0.057	0.096	-0.058

Standardized Within-Group Residuals:

	Min	Q1	Med	Q3
	-2.157276598	-0.389534045	0.007514936	0.366479340

Number of Observations: 324

We have no evidence of a long-term effect of duration of coma, only weak evidence of an effect at the end of 30 days.

Refitting PIQ with a similar model:

```
> fit.nlme.piq2 <-nlme( piq ~ b0 + b1*exp(-a*(dcoma-30)),
+   data = iq,
+   fixed = list( b0 ~ 1 + sqrt(dcoma) ,
+                 b1 ~ 1 + sqrt(dcoma) ,
+                 a ~ 1),
+   random = list( id = list( b0 ~ 1, b1~ 1),
+   start = list(
+     fixed = c(100, -.3, -10, 0,.1)),
+   control = list( maxIter = 100, returnObject = TRUE),
+   verbose = T,
+   subset = dcoma < 100)
```

which also converges quickly.

```
> summary( fit.nlme.piq2 )
```

```
. . . .  
Random effects:  
  Formula: list(b0 ~ 1, b1 ~ 1)  
  Level: id  
  Structure: General positive-definite, Log-Cholesky  
parametrization
```

	StdDev	Corr
b0.(Intercept)	1.303845e+01	b0.(I)
b1.(Intercept)	1.288991e-04	0.001
Residual	6.640711e+00	

```
Fixed effects: list(b0 ~ 1 + sqrt(dcoma), b1 ~ 1 +  
sqrt(dcoma), a ~ 1)
```

	Value	Std.Error	DF	t-value
b0.(Intercept)	98.45118	2.1594531	123	45.59079
b0.sqrt(dcoma)	-1.52286	0.5784392	123	-2.63270
b1.(Intercept)	-10.71808	2.6410052	123	-4.05833
b1.sqrt(dcoma)	-2.06639	0.8298037	123	-2.49022
a	0.00707	0.0014827	123	4.76892

These results contrast interestingly with VIQ.

**EXERCISE:** Note that **b1** has very small variability in  
What happens if you refit without a random effect for **b**

*Comparison of half-recovery times*

Half-recovery time =  $\frac{1}{\alpha \times \ln(2)}$

IQ	$\alpha$	half-recovery time
VIQ	0.02148	67 days
PIQ	0.00707	204 days

**EXERCISE:** Reparametrize the non-linear model form  
half-life instead of IRRR in the model. Are the results c

# *Visual comparisons of PIQ and VIQ*


#

```
windows( height = 7, width = 8.5)


pred <- expand.grid( dcoma = c(0,1,7,16,25,100),
                    dayspc = seq(30,365*2,5))
pred$piq <- predict( fit.nlme.piq2, pred, level = 0 )
pred$viq <- predict( fit.nlme.viq2, pred, level = 0 )

zz <- factor( paste( 'dcoma =', pred$dcoma))
pred$dcoma.lab <- reorder( zz, pred$dcoma)

td( col = c('blue','red'), lwd = 2)
xyplot( viq + piq ~ dayspc | dcoma.lab, pred, type =
        ylim = c(60,102),
        lwd = 2, auto.key = list(columns = 2, points = F,
```



```
pred2 <- expand.grid( dcoma = 0:100,  
                     dayspc = c(30,60,90,180,360,720))  
pred2$piq <- predict( fit.nlme.piq2, pred2, level = 0)  
pred2$viq <- predict( fit.nlme.viq2, pred2, level = 0)  
  
zz <- factor( paste( 'dayspc =', pred2$dayspc))  
pred2$dayspc.lab <- reorder( zz, pred2$dayspc)  
  
xyplot( viq + piq ~ dcoma | dayspc.lab, pred2, type =  
        ylim = c(60,102),  
        lwd = 2,  
        auto.key = list(columns = 2, points = F, lin
```




80



```
predviq <- expand.grid( dcoma = seq(0,100,10),  
                      dayspc = seq(30,720,30))  
predpiq <- predviq  
predviq $ iq <- predict( fit.nlme.viq2, predviq, level=1)  
predpiq $ iq <- predict( fit.nlme.piq2      , predpiq, level=1)  
predpiq$type <- factor( "PIQ" )  
predviq$type <- factor( "VIQ" )
```

```
wireframe( iq ~ dayspc + dcoma | type, Rbind( predpiq, predviq ) )
```

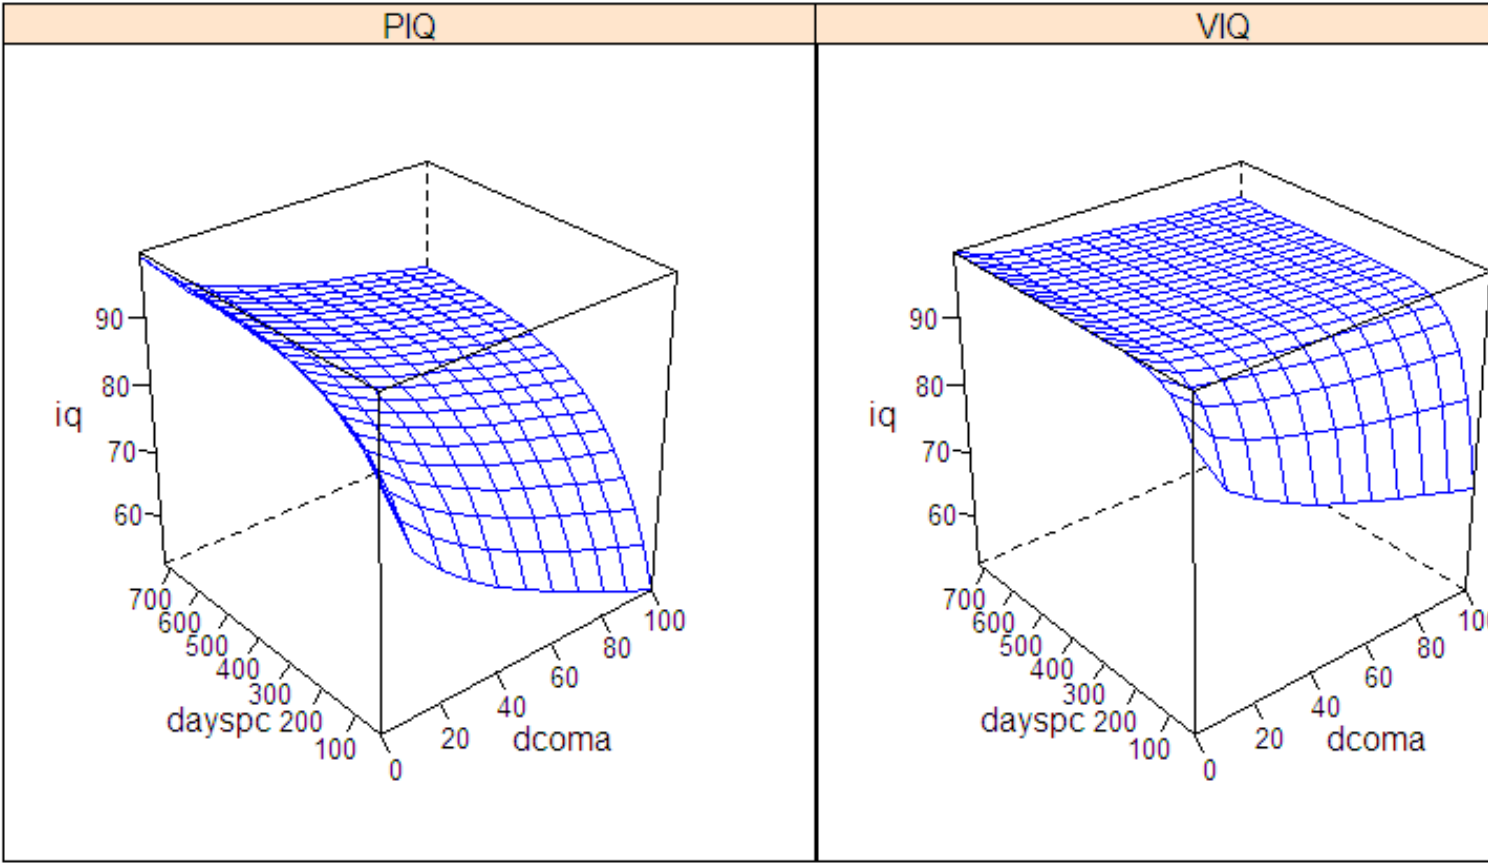
Default options and orientation for a wireframe plot:



```
wireframe( iq ~ dcoma + dayspc | type, Rbind( predpiq  
scales = list( arrows = F), col = 'blue')
```

Exchange axes for a more natural presentation of dayspc  
variable) and dcoma (Level 2)


Suppress arrows and get axis with values and tickmarks




```
wireframe( iq ~ dcoma + dayspc | type, Rbind(predpiq,  
  scales = list( arrows = F), col = 'blue',  
  xlab = 'Coma',  
  ylab = 'Days post coma',  
  screen = list( z = -65, x = -75 ))
```

With the 'screen' parameter, you can control the orientation of the graph. Here, the z axis is the vertical axis, the x axis, the axis in the surface of the screen and y, the horizontal axis pointing into the screen.

Rotation in the z axis of -65 degrees results in clockwise rotation of +65 degrees from the top and x-axis rotation of -75, tilted up by 75 degrees.



```
wireframe( iq ~ dcoma + dayspc , Rbind(predpiq, predv  
  groups = type,  
  scales = list( arrows = F), col = 'blue',  
  xlab = 'Coma',  
  ylab = 'Days post coma',  
  screen = list( z = -65, x = -75 ),  
  auto.key = list(columns=2, lines = T, points  
  alpha = .5)
```





See the Non-Linear Lab script for a multivariate model compares the parameters for PIQ and VIQ.